

Perceptrons and Structured Perceptrons

CS114B Lab 8

Kenneth Lai

March 26, 2021

Perceptrons

- ▶ Documents are characterized by features
 - ▶ No independence assumptions
- ▶ For each feature j :
 - ▶ Value x_j
 - ▶ Weight w_j
- ▶ Bias term b

- ▶ “Score” $z = \left(\sum_{j=1}^n w_j x_j \right) + b = \mathbf{w} \cdot \mathbf{x} + b$


Perceptrons

- ▶ Documents are characterized by features
 - ▶ No independence assumptions
- ▶ For each feature j :
 - ▶ Value x_j
 - ▶ Weight w_j
- ▶ Bias term b


- ▶ “Score” $z = \left(\sum_{j=1}^n w_j x_j \right) + b = \mathbf{w} \cdot \mathbf{x} + b$

- ▶ Does this look familiar?


Graphical Representation of Logistic Regression



Graphical Representation of a Neuron




Graphical Representation of a Perceptron




Graphical Representation of a Perceptron

- (Heaviside) step function H




Graphical Representation of a Perceptron



► (Heaviside) step function H

$$\text{► } H(z) = \begin{cases} 1, & \text{if } z > 0 \\ 0, & \text{if } z < 0 \end{cases}$$

Graphical Representation of a Perceptron




- (Heaviside) step function H

- $H(z) = \begin{cases} 1, & \text{if } z > 0 \\ 0, & \text{if } z < 0 \end{cases}$

- What if $z = 0$?

Graphical Representation of a Perceptron



- ▶ (Heaviside) step function H

- ▶ $H(z) = \begin{cases} 1, & \text{if } z > 0 \\ 0, & \text{if } z < 0 \end{cases}$

- ▶ What if $z = 0$?

- ▶ Set by convention
(1, 0, or 1/2)

Gradients in Perceptrons

- ▶ Hinge loss $L(\hat{y}, y) = (\hat{y} - y)z$

Gradients in Perceptrons

- ▶ Hinge loss $L(\hat{y}, y) = (\hat{y} - y)z$
- ▶ We want to compute $\frac{\partial L}{\partial w_j}$
- ▶ Chain Rule of calculus: $\frac{dy}{dx} = \frac{dy}{dz} \frac{dz}{dx}$
- ▶ Looking at the graph: $\frac{\partial L}{\partial w_j} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial z} \frac{\partial z}{\partial w_j}$

Gradients in Perceptrons

- ▶ Hinge loss $L(\hat{y}, y) = (\hat{y} - y)z$
- ▶ We want to compute $\frac{\partial L}{\partial w_j}$
- ▶ Chain Rule of calculus: $\frac{dy}{dx} = \frac{dy}{dz} \frac{dz}{dx}$
- ▶ Looking at the graph: $\frac{\partial L}{\partial w_j} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial z} \frac{\partial z}{\partial w_j}$
 - ▶ Problem: $\frac{\partial \hat{y}}{\partial z} = 0$ almost everywhere

Gradients in Perceptrons

- ▶ Hinge loss $L(\hat{y}, y) = (\hat{y} - y)z$
- ▶ We want to compute $\frac{\partial L}{\partial w_j}$
- ▶ Chain Rule of calculus: $\frac{dy}{dx} = \frac{dy}{dz} \frac{dz}{dx}$
- ▶ Looking at the graph: $\frac{\partial L}{\partial w_j} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial z} \frac{\partial z}{\partial w_j}$
 - ▶ Problem: $\frac{\partial \hat{y}}{\partial z} = 0$ almost everywhere
 - ▶ A change in z does not result in a change in \hat{y} unless you are at the threshold

Gradients in Perceptrons

- ▶ Hinge loss $L(\hat{y}, y) = (\hat{y} - y)z$
- ▶ We want to compute $\frac{\partial L}{\partial w_j}$
- ▶ Chain Rule of calculus: $\frac{dy}{dx} = \frac{dy}{dz} \frac{dz}{dx}$
- ▶ Looking at the graph: $\frac{\partial L}{\partial w_j} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial z} \frac{\partial z}{\partial w_j}$
 - ▶ Problem: $\frac{\partial \hat{y}}{\partial z} = 0$ almost everywhere
 - ▶ A change in z does not result in a change in \hat{y} unless you are at the threshold
 - ▶ Solution: consider $\frac{\partial L}{\partial z}$ directly

Gradients in Perceptrons

- ▶ Hinge loss $L(\hat{y}, y) = (\hat{y} - y)z$
- ▶ We want to compute $\frac{\partial L}{\partial w_j}$
- ▶ Chain Rule of calculus: $\frac{dy}{dx} = \frac{dy}{dz} \frac{dz}{dx}$
- ▶ Looking at the graph: $\frac{\partial L}{\partial w_j} = \frac{\partial L}{\partial z} \frac{\partial z}{\partial w_j}$

Gradients in Perceptrons

$$\blacktriangleright \frac{\partial L}{\partial w_j} = \frac{\partial L}{\partial z} \frac{\partial z}{\partial w_j}$$

Gradients in Perceptrons

- ▶ $\frac{\partial L}{\partial w_j} = \frac{\partial L}{\partial z} x_j$
 - ▶ $\frac{\partial z}{\partial w_j} = x_j$

Gradients in Perceptrons

- ▶ $\frac{\partial L}{\partial w_j} = (\hat{y} - y)x_j$
 - ▶ $\frac{\partial z}{\partial w_j} = x_j$
 - ▶ For the hinge loss: $\frac{\partial L}{\partial z} = \hat{y} - y$

Gradients in Perceptrons

- ▶ $\frac{\partial L}{\partial w_j} = (\hat{y} - y)x_j$
 - ▶ $\frac{\partial z}{\partial w_j} = x_j$
 - ▶ For the hinge loss: $\frac{\partial L}{\partial z} = \hat{y} - y$
- ▶ $\frac{\partial L}{\partial b} = \hat{y} - y$
 - ▶ $\frac{\partial z}{\partial b} = 1$
 - ▶ ...

Gradient Descent

- ▶ Initialize parameters $\theta = \mathbf{w}, b$ (randomly or $\mathbf{0}$)
- ▶ At each time step t :
 - ▶ Compute gradient ∇L
 - ▶ Move in direction of negative gradient
- ▶ $\theta_{t+1} = \theta_t - \eta \nabla L$

Gradient Descent

- ▶ Initialize parameters $\theta = \mathbf{w}, b$ (randomly or $\mathbf{0}$)
- ▶ At each time step t :
 - ▶ Compute gradient ∇L
 - ▶ Move in direction of negative gradient
- ▶ $\theta_{t+1} = \theta_t - \eta \nabla L$
- ▶ Does this look familiar?

Perceptron Learning Algorithm

- ▶ If $\hat{y} = y$, then $\hat{y} - y = 0$
 - ▶ Do nothing

Perceptron Learning Algorithm

- ▶ If $\hat{y} = y$, then $\hat{y} - y = 0$
 - ▶ Do nothing
- ▶ If $\hat{y} = 0$ and $y = 1$, then $\hat{y} - y = -1$
 - ▶ $\nabla L = -\mathbf{x}$
 - ▶ $\theta_{t+1} = \theta_t + \eta \mathbf{x}$
 - ▶ Increment weights

Perceptron Learning Algorithm

- ▶ If $\hat{y} = y$, then $\hat{y} - y = 0$
 - ▶ Do nothing
- ▶ If $\hat{y} = 0$ and $y = 1$, then $\hat{y} - y = -1$
 - ▶ $\nabla L = -\mathbf{x}$
 - ▶ $\theta_{t+1} = \theta_t + \eta \mathbf{x}$
 - ▶ Increment weights
- ▶ If $\hat{y} = 1$ and $y = 0$, then $\hat{y} - y = 1$
 - ▶ $\nabla L = \mathbf{x}$
 - ▶ $\theta_{t+1} = \theta_t - \eta \mathbf{x}$
 - ▶ Decrement weights

Multi-class Perceptrons

- ▶ Separate weights and bias terms for each class c

$$z_c = \left(\sum_{j=1}^n w_{jc} x_j \right) + b_c = \mathbf{w}_c \cdot \mathbf{x} + b_c$$

Multi-class Perceptrons

- ▶ Separate weights and bias terms for each class c

$$z_c = \left(\sum_{j=1}^n w_{jc} x_j \right) + b_c = \mathbf{w}_c \cdot \mathbf{x} + b_c$$

- ▶ $\hat{y} = \operatorname{argmax}_{k=1}^K z_k$

Multi-class Perceptrons

- ▶ Separate weights and bias terms for each class c

$$z_c = \left(\sum_{j=1}^n w_{jc} x_j \right) + b_c = \mathbf{w}_c \cdot \mathbf{x} + b_c$$

- ▶ $\hat{y} = \operatorname{argmax}_{k=1}^K z_k$
 - ▶ Unlike in multinomial logistic regression, \hat{y} does not have to be a vector

Multi-class Perceptrons

- ▶ Separate weights and bias terms for each class c

$$z_c = \left(\sum_{j=1}^n w_{jc} x_j \right) + b_c = \mathbf{w}_c \cdot \mathbf{x} + b_c$$

- ▶ $\hat{y} = \operatorname{argmax}_{k=1}^K z_k$
 - ▶ Unlike in multinomial logistic regression, \hat{y} does not have to be a vector
 - ▶ If we wanted to, we could define one-hot vectors $\hat{\mathbf{y}}$ and \mathbf{y}

Multi-class Perceptrons

- ▶ Separate weights and bias terms for each class c

$$z_c = \left(\sum_{j=1}^n w_{jc} x_j \right) + b_c = \mathbf{w}_c \cdot \mathbf{x} + b_c$$

- ▶ $\hat{y} = \operatorname{argmax}_{k=1}^K z_k$
 - ▶ Unlike in multinomial logistic regression, \hat{y} does not have to be a vector
 - ▶ If we wanted to, we could define one-hot vectors $\hat{\mathbf{y}}$ and \mathbf{y}
 - ▶ What if there is a tie?

Multi-class Perceptrons

- ▶ Separate weights and bias terms for each class c

$$z_c = \left(\sum_{j=1}^n w_{jc} x_j \right) + b_c = \mathbf{w}_c \cdot \mathbf{x} + b_c$$

- ▶ $\hat{y} = \operatorname{argmax}_{k=1}^K z_k$
 - ▶ Unlike in multinomial logistic regression, \hat{y} does not have to be a vector
 - ▶ If we wanted to, we could define one-hot vectors $\hat{\mathbf{y}}$ and \mathbf{y}
 - ▶ What if there is a tie?
 - ▶ Do whatever `numpy.argmax` does

Multi-class Perceptrons

- ▶ Separate weights and bias terms for each class c

$$z_c = \left(\sum_{j=1}^n w_{jc} x_j \right) + b_c = \mathbf{w}_c \cdot \mathbf{x} + b_c$$

- ▶ $\hat{y} = \operatorname{argmax}_{k=1}^K z_k$
 - ▶ Unlike in multinomial logistic regression, \hat{y} does not have to be a vector
 - ▶ If we wanted to, we could define one-hot vectors $\hat{\mathbf{y}}$ and \mathbf{y}
 - ▶ What if there is a tie?
 - ▶ Do whatever `numpy.argmax` does
- ▶ Hinge loss $L(\hat{y}, y) = z_{\hat{y}} - z_y$

Multi-class Perceptron Learning Algorithm

- ▶ If $\hat{y} = y$, then do nothing
- ▶ Else:

Multi-class Perceptron Learning Algorithm

- ▶ If $\hat{y} = y$, then do nothing
- ▶ Else:
 - ▶ For the correct class y , $\frac{\partial L}{\partial z_y} = -1$
 - ▶ $(\nabla L)_y = -\mathbf{x}$
 - ▶ $(\theta_y)_{t+1} = (\theta_y)_t + \eta \mathbf{x}$
 - ▶ Increment weights

Multi-class Perceptron Learning Algorithm

- ▶ If $\hat{y} = y$, then do nothing
- ▶ Else:
 - ▶ For the correct class y , $\frac{\partial L}{\partial z_y} = -1$
 - ▶ $(\nabla L)_y = -\mathbf{x}$
 - ▶ $(\theta_y)_{t+1} = (\theta_y)_t + \eta \mathbf{x}$
 - ▶ Increment weights
 - ▶ For the predicted class \hat{y} , $\frac{\partial L}{\partial z_{\hat{y}}} = 1$
 - ▶ $(\nabla L)_{\hat{y}} = \mathbf{x}$
 - ▶ $(\theta_{\hat{y}})_{t+1} = (\theta_{\hat{y}})_t - \eta \mathbf{x}$
 - ▶ Decrement weights

Multi-class Perceptron Learning Algorithm

- ▶ If $\hat{y} = y$, then do nothing
- ▶ Else:
 - ▶ For the correct class y , $\frac{\partial L}{\partial z_y} = -1$
 - ▶ $(\nabla L)_y = -\mathbf{x}$
 - ▶ $(\theta_y)_{t+1} = (\theta_y)_t + \eta \mathbf{x}$
 - ▶ Increment weights
 - ▶ For the predicted class \hat{y} , $\frac{\partial L}{\partial z_{\hat{y}}} = 1$
 - ▶ $(\nabla L)_{\hat{y}} = \mathbf{x}$
 - ▶ $(\theta_{\hat{y}})_{t+1} = (\theta_{\hat{y}})_t - \eta \mathbf{x}$
 - ▶ Decrement weights
 - ▶ For other classes, do nothing

Structured Perceptrons

- ▶ Perceptrons for sequence labeling

Structured Perceptrons

- ▶ Perceptrons for sequence labeling
 - ▶ Given a sequence $X = [x_1, \dots, x_T]$, predict labels $Y = [y_1, \dots, y_T]$

Structured Perceptrons

- ▶ Perceptrons for sequence labeling
 - ▶ Given a sequence $X = [x_1, \dots, x_T]$, predict labels $Y = [y_1, \dots, y_T]$
 - ▶ We do not care about the probability $P(Y|X)$, just which Y has the highest score Z

Structured Perceptrons

- ▶ Perceptrons for sequence labeling
 - ▶ Given a sequence $X = [x_1, \dots, x_T]$, predict labels $Y = [y_1, \dots, y_T]$
 - ▶ We do not care about the probability $P(Y|X)$, just which Y has the highest score Z
- ▶ $\hat{Y} = \operatorname{argmax}_{k \in K^T} Z_k$

Structured Perceptrons

- ▶ Score Z decomposes into a sum of local parts

Structured Perceptrons

- ▶ Score Z decomposes into a sum of local parts
 - ▶ At each time step i , for each possible combination of current tag y_i and previous tag y_{i-1} , compute a local score $z(y_i, y_{i-1})$

Structured Perceptrons

- ▶ Score Z decomposes into a sum of local parts
 - ▶ At each time step i , for each possible combination of current tag y_i and previous tag y_{i-1} , compute a local score $z(y_i, y_{i-1})$
 - ▶ Use the Viterbi algorithm to combine the local scores across the sequence, and find the argmax

Structured Perceptrons

- ▶ Suppose that at each time step i , we want to predict the current tag y_i using the following features:

Structured Perceptrons

- ▶ Suppose that at each time step i , we want to predict the current tag y_i using the following features:
 - ▶ Previous tag y_{i-1}

Structured Perceptrons

- ▶ Suppose that at each time step i , we want to predict the current tag y_i using the following features:
 - ▶ Previous tag y_{i-1}
 - ▶ At the beginning of the sentence, let the previous tag y_{i-1} be the start symbol $\langle S \rangle$

Structured Perceptrons

- ▶ Suppose that at each time step i , we want to predict the current tag y_i using the following features:
 - ▶ Previous tag y_{i-1}
 - ▶ At the beginning of the sentence, let the previous tag y_{i-1} be the start symbol $\langle S \rangle$
 - ▶ Current word x_i

Structured Perceptrons

- ▶ Suppose that at each time step i , we want to predict the current tag y_i using the following features:
 - ▶ Previous tag y_{i-1}
 - ▶ At the beginning of the sentence, let the previous tag y_{i-1} be the start symbol $\langle S \rangle$
 - ▶ Current word x_i
- ▶ For simplicity, we will assume that these are the only features, and we will ignore the bias term

Structured Perceptrons

- ▶ Suppose that at each time step i , we want to predict the current tag y_i using the following features:
 - ▶ Previous tag y_{i-1}
 - ▶ At the beginning of the sentence, let the previous tag y_{i-1} be the start symbol $\langle S \rangle$
 - ▶ Current word x_i
- ▶ For simplicity, we will assume that these are the only features, and we will ignore the bias term
- ▶ Let $\mathbf{f}(X, y_i, y_{i-1}, i)$ be the feature vector at time step i

Structured Perceptrons

- ▶ Suppose that at each time step i , we want to predict the current tag y_i using the following features:
 - ▶ Previous tag y_{i-1}
 - ▶ At the beginning of the sentence, let the previous tag y_{i-1} be the start symbol $\langle S \rangle$
 - ▶ Current word x_i
- ▶ For simplicity, we will assume that these are the only features, and we will ignore the bias term
- ▶ Let $\mathbf{f}(X, y_i, y_{i-1}, i)$ be the feature vector at time step i
 - ▶ Using \mathbf{f} instead of \mathbf{x} , because features can include more than just the input

Structured Perceptrons

- ▶ We can arrange our weight matrix Θ as follows:

$$\begin{bmatrix} \text{---} \\ \text{---} \end{bmatrix}$$

Structured Perceptrons

- ▶ We can arrange our weight matrix Θ as follows:

$$\begin{bmatrix} \pi \\ \hline \\ \hline \\ \end{bmatrix}$$

- ▶ Initial features
 - ▶ $y_{i-1} = \langle S \rangle, y_i = \dots$

Structured Perceptrons

- ▶ We can arrange our weight matrix Θ as follows:

$$\begin{bmatrix} \pi \\ \hline \mathbf{A} \\ \hline \end{bmatrix}$$

- ▶ Initial features
 - ▶ $y_{i-1} = \langle S \rangle, y_i = \dots$
- ▶ Transition features
 - ▶ $y_{i-1} = \dots, y_i = \dots$

Structured Perceptrons

- ▶ We can arrange our weight matrix Θ as follows:

$$\left[\begin{array}{c} \pi \\ \hline \mathbf{A} \\ \hline \mathbf{B} \end{array} \right]$$

- ▶ Initial features
 - ▶ $y_{i-1} = \langle S \rangle, y_i = \dots$
- ▶ Transition features
 - ▶ $y_{i-1} = \dots, y_i = \dots$
- ▶ Emission features
 - ▶ $x_i = \dots, y_i = \dots$

Initialization Step

- ▶ We want to compute local scores $z(y_1, \langle S \rangle)$ for each possible y_1

Initialization Step

- ▶ We want to compute local scores $z(y_1, \langle S \rangle)$ for each possible y_1
 - ▶ These are the elements of $\mathbf{z}_1 = \mathbf{f}(X, y_1, \langle S \rangle, 1) \cdot \Theta$

Initialization Step

$$\mathbf{z}_1 = \mathbf{f}(X, y_1, \langle S \rangle, 1) \cdot \Theta$$

$$= \begin{bmatrix} & | & & | \end{bmatrix} \cdot \begin{bmatrix} \pi \\ \hline \mathbf{A} \\ \hline \mathbf{B} \end{bmatrix}$$

Initialization Step

$$\mathbf{z}_1 = \mathbf{f}(X, y_1, \langle S \rangle, 1) \cdot \Theta$$

$$= \left[\begin{array}{c|c} 1 & \end{array} \right]$$

$$\cdot \left[\begin{array}{c} \pi \\ \hline \mathbf{A} \\ \hline \mathbf{B} \end{array} \right]$$

- We know that $y_{i-1} = \langle S \rangle$

Initialization Step

$$\mathbf{z}_1 = \mathbf{f}(X, y_1, \langle S \rangle, 1) \cdot \Theta$$

$$= \left[\begin{array}{c|c|c} 1 & \mathbf{0} & \end{array} \right] \cdot \left[\begin{array}{c} \pi \\ \hline \mathbf{A} \\ \hline \mathbf{B} \end{array} \right]$$

- We know that y_{i-1} cannot be any other tag

Initialization Step

$$\mathbf{z}_1 = \mathbf{f}(X, y_1, \langle S \rangle, 1) \cdot \Theta$$
$$= \left[\begin{array}{c|c|c} 1 & \mathbf{0} & \mathbf{1}_{\{x_1 = o_1\}} \end{array} \right] \cdot \left[\begin{array}{c} \pi \\ \hline \mathbf{A} \\ \hline \mathbf{B} \end{array} \right]$$

- One-hot vector of the first word

Initialization Step

$$\mathbf{z}_1 = \mathbf{f}(X, y_1, \langle S \rangle, 1) \cdot \Theta$$

$$= \begin{bmatrix} 1 & | & \mathbf{0} & | & \mathbf{1}_{\{x_1 = o_1\}} \end{bmatrix} \cdot \begin{bmatrix} \pi \\ \hline \mathbf{A} \\ \hline \mathbf{B} \end{bmatrix}$$

$$= 1 \cdot \pi + \mathbf{0} \cdot \mathbf{A} + \mathbf{1}_{\{x_1 = o_1\}} \cdot \mathbf{B}$$

Initialization Step

$$\mathbf{z}_1 = \mathbf{f}(X, y_1, \langle S \rangle, 1) \cdot \Theta$$

$$= \begin{bmatrix} 1 & | & \mathbf{0} & | & \mathbf{1}_{\{x_1 = o_1\}} \end{bmatrix} \cdot \begin{bmatrix} \pi \\ \hline \mathbf{A} \\ \hline \mathbf{B} \end{bmatrix}$$

$$= 1 \cdot \pi + \mathbf{0} \cdot \mathbf{A} + \mathbf{1}_{\{x_1 = o_1\}} \cdot \mathbf{B}$$

$$= \pi + \mathbf{b}(o_1)$$

Initialization Step

$$\mathbf{z}_1 = \mathbf{f}(X, y_1, \langle S \rangle, 1) \cdot \Theta$$

$$= \begin{bmatrix} 1 & | & \mathbf{0} & | & \mathbf{1}\{x_1 = o_1\} \end{bmatrix} \cdot \begin{bmatrix} \pi \\ \hline \mathbf{A} \\ \hline \mathbf{B} \end{bmatrix}$$

$$= 1 \cdot \pi + \mathbf{0} \cdot \mathbf{A} + \mathbf{1}\{x_1 = o_1\} \cdot \mathbf{B}$$

$$= \pi + \mathbf{b}(o_1)$$

- These local scores go into the first column of the Viterbi trellis

Recursion Step

- ▶ We want to compute local scores $z(y_i, y_{i-1})$ for each possible combination of y_i and y_{i-1}

Recursion Step

- ▶ We want to compute local scores $z(y_i, y_{i-1})$ for each possible combination of y_i and y_{i-1}
 - ▶ We can stack the feature vectors for each possible y_{i-1} , $\mathbf{f}(X, y_i, y_{i-1}, i)$, on top of each other, in order

Recursion Step

- ▶ We want to compute local scores $z(y_i, y_{i-1})$ for each possible combination of y_i and y_{i-1}
 - ▶ We can stack the feature vectors for each possible y_{i-1} , $\mathbf{f}(X, y_i, y_{i-1}, i)$, on top of each other, in order
 - ▶ Form a feature matrix \mathbf{F}_i

Recursion Step

- ▶ We want to compute local scores $z(y_i, y_{i-1})$ for each possible combination of y_i and y_{i-1}
 - ▶ We can stack the feature vectors for each possible y_{i-1} , $\mathbf{f}(X, y_i, y_{i-1}, i)$, on top of each other, in order
 - ▶ Form a feature matrix \mathbf{F}_i
 - ▶ Compute $\mathbf{Z}_i = \mathbf{F}_i \cdot \Theta$

Recursion Step

$$\mathbf{Z}_i = \mathbf{F}_i \cdot \boldsymbol{\Theta}$$

$$= \begin{bmatrix} | & | & | \end{bmatrix} \cdot \begin{bmatrix} \pi \\ \hline \mathbf{A} \\ \hline \mathbf{B} \end{bmatrix}$$

Recursion Step

$$\mathbf{Z}_i = \mathbf{F}_i \cdot \Theta$$

$$= \left[\begin{array}{c|c} \mathbf{0} & \end{array} \right] \cdot \left[\begin{array}{c} \pi \\ \hline \mathbf{A} \\ \hline \mathbf{B} \end{array} \right]$$

- We know that $y_{i-1} \neq \langle S \rangle$

Recursion Step

$$\mathbf{Z}_i = \mathbf{F}_i \cdot \Theta$$

$$= \left[\begin{array}{c|c} \mathbf{0} & \mathbf{I} \end{array} \right] \cdot \left[\begin{array}{c} \pi \\ \hline \mathbf{A} \\ \hline \mathbf{B} \end{array} \right]$$

- Identity matrix!

Recursion Step

$$\mathbf{Z}_i = \mathbf{F}_i \cdot \Theta$$

$$= \left[\begin{array}{c|c|c} \mathbf{0} & \mathbf{I} & \mathbf{1}_{\{x_i = o_i\}} \end{array} \right] \cdot \left[\begin{array}{c} \pi \\ \hline \mathbf{A} \\ \hline \mathbf{B} \end{array} \right]$$

- Stack of one-hot vectors

Recursion Step

$$\mathbf{Z}_i = \mathbf{F}_i \cdot \boldsymbol{\Theta}$$

$$= \left[\begin{array}{c|c|c} \mathbf{0} & \mathbf{I} & \mathbf{1}_{\{x_i = o_i\}} \end{array} \right] \cdot \left[\begin{array}{c} \pi \\ \hline \mathbf{A} \\ \hline \mathbf{B} \end{array} \right]$$
$$= 0 \cdot \pi + \mathbf{I} \cdot \mathbf{A} + \mathbf{1}_{\{x_i = o_i\}} \cdot \mathbf{B}$$

Recursion Step

$$\mathbf{Z}_i = \mathbf{F}_i \cdot \Theta$$

$$= \left[\begin{array}{c|c|c} \mathbf{0} & \mathbf{I} & \mathbf{1}_{\{x_i = o_i\}} \end{array} \right] \cdot \left[\begin{array}{c} \pi \\ \hline \mathbf{A} \\ \hline \mathbf{B} \end{array} \right]$$

$$= \mathbf{0} \cdot \pi + \mathbf{I} \cdot \mathbf{A} + \mathbf{1}_{\{x_i = o_i\}} \cdot \mathbf{B}$$

$$= \mathbf{A} + \mathbf{b}(o_i)$$

Recursion Step

$$\mathbf{Z}_i = \mathbf{F}_i \cdot \Theta$$

$$= \left[\begin{array}{c|c|c} \mathbf{0} & \mathbf{I} & \mathbf{1}_{\{x_i = o_i\}} \end{array} \right] \cdot \left[\begin{array}{c} \pi \\ \hline \mathbf{A} \\ \hline \mathbf{B} \end{array} \right]$$

$$= \mathbf{0} \cdot \pi + \mathbf{I} \cdot \mathbf{A} + \mathbf{1}_{\{x_i = o_i\}} \cdot \mathbf{B}$$

$$= \mathbf{A} + \mathbf{b}(o_i)$$

- Use the Viterbi algorithm to combine these local scores with scores from the rest of the sequence

Structured Perceptron Learning Algorithm

- ▶ Use the Viterbi algorithm to compute the best tag sequence

Structured Perceptron Learning Algorithm

- ▶ Use the Viterbi algorithm to compute the best tag sequence
- ▶ If $\hat{Y} = Y$, then do nothing

Structured Perceptron Learning Algorithm

- ▶ Use the Viterbi algorithm to compute the best tag sequence
- ▶ If $\hat{Y} = Y$, then do nothing
- ▶ Else:
 - ▶ Increment weights for features in Y , decrement weights for features in \hat{Y}

Structured Perceptron Learning Algorithm

- ▶ Use the Viterbi algorithm to compute the best tag sequence
- ▶ If $\hat{Y} = Y$, then do nothing
- ▶ Else:
 - ▶ Increment weights for features in Y , decrement weights for features in \hat{Y}
 - ▶ In other words, for each time step i :

Structured Perceptron Learning Algorithm

- ▶ Use the Viterbi algorithm to compute the best tag sequence
- ▶ If $\hat{Y} = Y$, then do nothing
- ▶ Else:
 - ▶ Increment weights for features in Y , decrement weights for features in \hat{Y}
 - ▶ In other words, for each time step i :
 - ▶ Increment weights for features in y_i , decrement weights for features in \hat{y}_i

Structured Perceptron Learning Algorithm

- ▶ Use the Viterbi algorithm to compute the best tag sequence
- ▶ If $\hat{Y} = Y$, then do nothing
- ▶ Else:
 - ▶ Increment weights for features in Y , decrement weights for features in \hat{Y}
 - ▶ In other words, for each time step i :
 - ▶ Increment weights for features in y_i , decrement weights for features in \hat{y}_i
 - ▶ Nothing fancy; no Numpy tricks needed